

Computer Vision - Lecture 7

Segmentation as Energy Minimization

16.11.2016

Bastian Leibe

RWTH Aachen

<http://www.vision.rwth-aachen.de>

leibe@vision.rwth-aachen.de

Announcements


- Please don't forget to register for the exam!
 - On the Campus system

Course Outline

- Image Processing Basics
- Segmentation
 - Segmentation and Grouping
 - Segmentation as Energy Minimization
- Recognition
 - Global Representations
 - Subspace representations
- Local Features & Matching
- Object Categorization
- 3D Reconstruction
- Motion and Tracking

Recap: Image Segmentation

- Goal: identify groups of pixels that go together




Recap: K-Means Clustering

- Basic idea: randomly initialize the k cluster centers, and iterate between the two steps we just saw.
 1. Randomly initialize the cluster centers, c_1, \dots, c_k
 2. Given cluster centers, determine points in each cluster
 - For each point p , find the closest c_i . Put p into cluster i
 3. Given points in each cluster, solve for c_i
 - Set c_i to be the mean of points in cluster i
 4. If c_i have changed, repeat Step 2
- Properties
 - Will always converge to *some* solution
 - Can be a “local minimum”
 - Does not always find the global minimum of objective function:

$$\sum_{\text{clusters } i} \sum_{\text{points } p \text{ in cluster } i} \|p - c_i\|^2$$



Recap: Expectation Maximization (EM)



- **Goal**

- Find blob parameters θ that maximize the likelihood function:

$$p(\text{data}|\theta) = \prod_{n=1}^N p(\mathbf{x}_n|\theta)$$

- **Approach:**

1. E-step: given current guess of blobs, compute ownership of each point
2. M-step: given ownership probabilities, update blobs to maximize likelihood function
3. Repeat until convergence

Recap: EM Algorithm

- Expectation-Maximization (EM) Algorithm

- **E-Step:** softly assign samples to mixture components

$$\gamma_j(\mathbf{x}_n) \leftarrow \frac{\pi_j \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x}_n | \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)} \quad \forall j = 1, \dots, K, \quad n = 1, \dots, N$$

- **M-Step:** re-estimate the parameters (separately for each mixture component) based on the soft assignments

$$\hat{N}_j \leftarrow \sum_{n=1}^N \gamma_j(\mathbf{x}_n) = \text{soft number of samples labeled } j$$

$$\hat{\pi}_j^{\text{new}} \leftarrow \frac{\hat{N}_j}{N}$$


$$\hat{\boldsymbol{\mu}}_j^{\text{new}} \leftarrow \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(\mathbf{x}_n) \mathbf{x}_n$$

$$\hat{\boldsymbol{\Sigma}}_j^{\text{new}} \leftarrow \frac{1}{\hat{N}_j} \sum_{n=1}^N \gamma_j(\mathbf{x}_n) (\mathbf{x}_n - \hat{\boldsymbol{\mu}}_j^{\text{new}})(\mathbf{x}_n - \hat{\boldsymbol{\mu}}_j^{\text{new}})^T$$


MoG Color Models for Image Segmentation



(a) input image



(b) user input




(c) inferred segmentation

- **User assisted image segmentation**

- User marks two regions for foreground and background.
 - Learn a MoG model for the color values in each region.
 - Use those models to classify all other pixels.
- ⇒ Simple segmentation procedure
(building block for more complex applications)

Recap: Mean-Shift Algorithm




- **Iterative Mode Search**

1. Initialize random seed, and window W
2. Calculate center of gravity (the “mean”) of W : $\sum_{x \in W} x H(x)$
3. Shift the search window to the mean
4. Repeat Step 2 until convergence



Recap: Mean-Shift Clustering

- Cluster: all data points in the attraction basin of a mode
- Attraction basin: the region for which all trajectories lead to the same mode




Recap: Mean-Shift Segmentation

- Find features (color, gradients, texture, etc)
- Initialize windows at individual pixel locations
- Perform mean shift for each window until convergence
- Merge windows that end up near the same “peak” or mode



Back to the Image Segmentation Problem...


- Goal: identify groups of pixels that go together



- Up to now, we have focused on ways to group pixels into image segments based on their appearance...
 - Segmentation as clustering.
- We also want to enforce region constraints.
 - Spatial consistency
 - Smooth borders


Topics of This Lecture

- **Segmentation as Energy Minimization**
 - Markov Random Fields
 - Energy formulation
- **Graph cuts for image segmentation**
 - Basic idea
 - s-t Mincut algorithm
 - Extension to non-binary case
- **Applications**
 - Interactive segmentation



Markov Random Fields


- Allow rich probabilistic models for images
- But built in a local, modular way
 - Learn local effects, get global effects out



MRF Nodes as Pixels




Original image



Degraded image



Reconstruction
from MRF modeling
pixel neighborhood
statistics



Network Joint Probability

$$p(\mathbf{x}, \mathbf{y}) = \prod_i \Phi(x_i, y_i) \prod_{i,j} \Psi(x_i, x_j)$$

Diagram illustrating the components of the network joint probability equation:

- \mathbf{x} (Scene) and \mathbf{y} (Image) are inputs to the first product.
- $\Phi(x_i, y_i)$ is the **Image-scene compatibility function**.
- i represents **Local observations**.
- $\Psi(x_i, x_j)$ is the **Scene-scene compatibility function**.
- i, j represent **Neighboring scene nodes**.

Energy Formulation

- Joint probability

$$p(\mathbf{x}, \mathbf{y}) = \prod_i \Phi(x_i, y_i) \prod_{i,j} \Psi(x_i, x_j)$$

- Maximizing the joint probability is the same as minimizing the negative log

$$-\log p(\mathbf{x}, \mathbf{y}) = -\sum_i \log \Phi(x_i, y_i) - \sum_{i,j} \log \Psi(x_i, x_j)$$


$$E(\mathbf{x}, \mathbf{y}) = \sum_i \phi(x_i, y_i) + \sum_{i,j} \psi(x_i, x_j)$$

- This is similar to free-energy problems in statistical mechanics (spin glass theory). We therefore draw the analogy and call E an *energy function*.
- ϕ and ψ are called *potentials*.

Energy Formulation

- Energy function

$$E(\mathbf{x}, \mathbf{y}) = \sum_i \underbrace{\phi(x_i, y_i)}_{\text{Single-node potentials}} + \sum_{i,j} \underbrace{\psi(x_i, x_j)}_{\text{Pairwise potentials}}$$



- Single-node potentials ϕ (“unary potentials”)


- Encode local information about the given pixel/patch
- How likely is a pixel/patch to belong to a certain class (e.g. foreground/background)?

- Pairwise potentials ψ

- Encode neighborhood information
- How different is a pixel/patch’s label from that of its neighbor? (e.g. based on intensity/color/texture difference, edges)

Energy Minimization


- **Goal:**
 - Infer the optimal labeling of the MRF.
- Many inference algorithms are available, e.g.
 - Gibbs sampling, simulated annealing
 - Iterated conditional modes (ICM)
 - Variational methods
 - Belief propagation
 - **Graph cuts**
- Recently, Graph Cuts have become a popular tool
 - **Only suitable for a certain class of energy functions**
 - But the solution can be obtained very fast for typical vision problems (~1MPixel/sec).



see lecture
Machine Learning!

Topics of This Lecture




- Segmentation as Energy Minimization
 - Markov Random Fields
 - Energy formulation
- Graph cuts for image segmentation
 - Basic idea
 - s-t Mincut algorithm
 - Extension to non-binary case
- Applications
 - Interactive segmentation




Binary case

Graph Cuts for Optimal Boundary Detection

- Idea: convert MRF into source-sink graph



Minimum cost cut can be
computed in polynomial time
(max-flow/min-cut algorithms)





Simple Example of Energy

$$E(\mathbf{x}, \mathbf{y}) = \sum_i \phi_i(x_i) + \sum_{i,j} w_{ij} \cdot \delta(x_i \neq x_j)$$

Unary terms



Pairwise terms



$$x \in \{s, t\}$$

(binary object segmentation)

Adding Regional Properties



Regional bias example

Suppose I^s and I^t are given
“expected” intensities
of **object** and **background**





$$\phi_i(s) \propto \exp \left(- \| I_i - I^s \|^2 / 2\sigma^2 \right)$$

$$\phi_i(t) \propto \exp \left(- \| I_i - I^t \|^2 / 2\sigma^2 \right)$$

NOTE: hard constraints are not required, in general.

Adding Regional Properties



“expected” intensities of
object and **background**
 I^s and I^t
 can be re-estimated




$$\phi_i(s) \propto \exp \left(- \| I_i - I^s \|^2 / 2\sigma^2 \right)$$

$$\phi_i(t) \propto \exp \left(- \| I_i - I^t \|^2 / 2\sigma^2 \right)$$


EM-style optimization

Adding Regional Properties

- More generally, regional bias can be based on any intensity models of object and background



$$\phi_i(L_i) = -\log p(I_i | L_i)$$



given object and background intensity
histograms

How to Set the Potentials? Some Examples

- Color potentials

- e.g., modeled with a Mixture of Gaussians

$$\phi(x_i, y_i; \theta_\phi) = \log \sum_k \theta_\phi(x_i, k) p(k|x_i) \mathcal{N}(y_i; \mu_k, \Sigma_k)$$

- Edge potentials

- E.g., a “contrast sensitive Potts model”

$$\psi(x_i, x_j, g_{ij}(\mathbf{y}); \theta_\psi) = -\theta_\psi g_{ij}(\mathbf{y}) \delta(x_i \neq x_j)$$

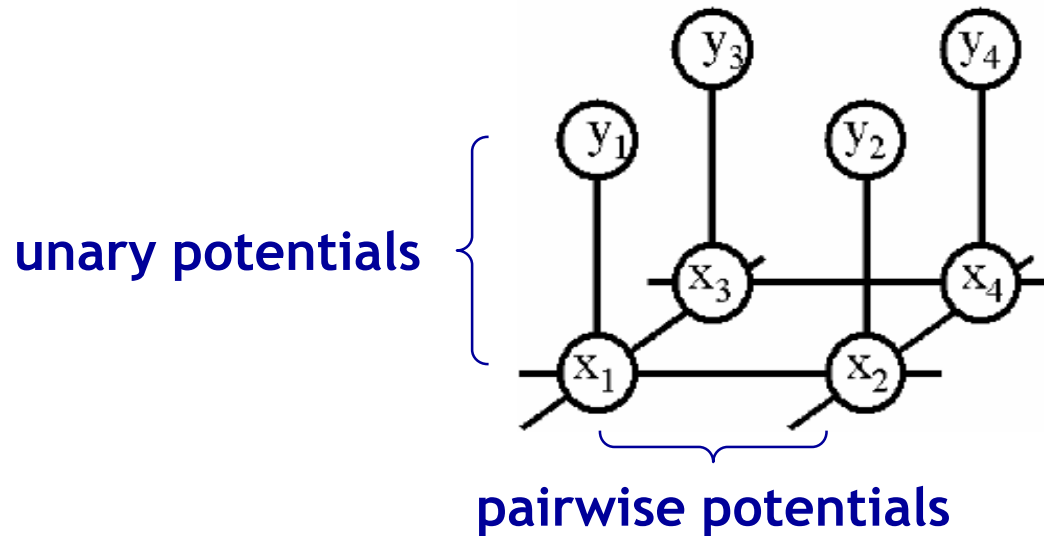
where

$$g_{ij}(\mathbf{y}) = e^{-\beta \|y_i - y_j\|^2} \quad \beta = \frac{1}{2} (\text{avg} (\|y_i - y_j\|^2))^{-1}$$

- Parameters θ_ϕ , θ_ψ need to be learned, too!

Example: MRF for Image Segmentation


- MRF structure



Data (D)



Unary likelihood




Pair-wise Terms




MAP Solution

Topics of This Lecture

- Segmentation as Energy Minimization
 - Markov Random Fields
 - Energy formulation
- Graph cuts for image segmentation
 - Basic idea
 - **s-t Mincut algorithm**
 - Extension to non-binary case
- Applications
 - Interactive segmentation



How Does it Work? The s-t-Mincut Problem



Graph (V, E, C)

Vertices $V = \{v_1, v_2 \dots v_n\}$

Edges $E = \{(v_1, v_2) \dots\}$

Costs $C = \{c_{(1, 2)} \dots\}$


The s-t-Mincut Problem

What is an st-cut?

An st-cut (S, T) divides the nodes between source and sink.

What is the cost of a st-cut?

Sum of cost of all edges going from S to T



$$5 + 2 + 9 = 16$$

The s-t-Mincut Problem

What is an st-cut?


An st-cut (S, T) divides the nodes between source and sink.

What is the cost of a st-cut?

Sum of cost of all edges going from S to T

What is the st-mincut?

st-cut with the minimum cost




$$2 + 1 + 4 = 7$$

How to Compute the s-t-Mincut?

Solve the dual maximum flow problem

Compute the maximum flow
between Source and Sink



Constraints

Edges: $\text{Flow} < \text{Capacity}$

Nodes: $\text{Flow in} = \text{Flow out}$

Min-cut/Max-flow Theorem

In every network, the maximum flow
equals the cost of the st-mincut

History of Maxflow Algorithms

Augmenting Path and Push-Relabel

year	discoverer(s)	bound
1951	Dantzig	$O(n^2 m U)$
1955	Ford & Fulkerson	$O(m^2 U)$
1970	Dinitz	$O(n^2 m)$
1972	Edmonds & Karp	$O(m^2 \log U)$
1973	Dinitz	$O(nm \log U)$
1974	Karzanov	$O(n^3)$
1977	Cherkassky	$O(n^2 m^{1/2})$
1980	Galil & Naamad	$O(nm \log^2 n)$
1983	Sleator & Tarjan	$O(nm \log n)$
1986	Goldberg & Tarjan	$O(nm \log(n^2/m))$
1987	Ahuja & Orlin	$O(nm + n^2 \log U)$
1987	Ahuja et al.	$O(nm \log(n \sqrt{\log U/m}))$
1989	Cheriyani & Hagerup	$E(nm + n^2 \log^2 n)$
1990	Cheriyani et al.	$O(n^3 / \log n)$
1990	Alon	$O(nm + n^{8/3} \log n)$
1992	King et al.	$O(nm + n^{2+\epsilon})$
1993	Phillips & Westbrook	$O(nm(\log_{m/n} n + \log^{2+\epsilon} n))$
1994	King et al.	$O(nm \log_{m/(n \log n)} n)$
1997	Goldberg & Rao	$O(m^{3/2} \log(n^2/m) \log U)$ $O(n^{2/3} m \log(n^2/m) \log U)$

n : #nodes


m : #edges

U : maximum
edge weight

**Algorithms
assume non-
negative edge
weights**

Maxflow Algorithms

Flow = 0




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 0




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 0 + 2




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 2 + 4




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 6




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 6




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 6 + 1




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 7




Augmenting Path Based Algorithms

1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Maxflow Algorithms

Flow = 7



Augmenting Path Based Algorithms


1. Find path from source to sink with positive capacity
2. Push maximum possible flow through this path
3. Repeat until no path can be found

Algorithms assume non-negative capacity

Applications: Maxflow in Computer Vision

- Specialized algorithms for vision problems


- Grid graphs
- Low connectivity ($m \sim O(n)$)



- Dual search tree augmenting path algorithm

[Boykov and Kolmogorov PAMI 2004]

- Finds approximate shortest augmenting paths efficiently.
- High worst-case time complexity.
- Empirically outperforms other algorithms on vision problems.
- Efficient code available on the web




<http://www.cs.ucl.ac.uk/staff/V.Kolmogorov/software.html>

When Can s-t Graph Cuts Be Applied?

$$E(L) = \sum_{\substack{p \\ \text{t-links}}} \overset{\text{Unary potentials}}{E_p(L_p)} + \sum_{\substack{pq \in N \\ \text{n-links}}} \overset{\text{Pairwise potentials}}{E(L_p, L_q)} \quad L_p \in \{s, t\}$$

- s-t graph cuts can only globally minimize **binary energies that are submodular**. [Boros & Hummer, 2002, Kolmogorov & Zabih, 2004]

$E(L)$ can be minimized
by s-t graph cuts




$$E(s, s) + E(t, t) \leq E(s, t) + E(t, s)$$

Submodularity (“convexity”)

- Submodularity is the discrete equivalent to convexity.
 - Implies that every local energy minimum is a global minimum.
 - \Rightarrow Solution will be globally optimal.

Topics of This Lecture

- Segmentation as Energy Minimization
 - Markov Random Fields
 - Energy formulation
- Graph cuts for image segmentation
 - Basic idea
 - s-t Mincut algorithm
 - Extension to non-binary case
- Applications
 - Interactive segmentation




Dealing with Non-Binary Cases

- Limitation to binary energies is often a nuisance.
⇒ E.g. binary segmentation only...
- We would like to solve also multi-label problems.
 - The bad news: Problem is NP-hard with 3 or more labels!
- There exist some approximation algorithms which extend graph cuts to the multi-label case:
 - α -Expansion
 - $\alpha\beta$ -Swap
- They are no longer guaranteed to return the globally optimal result.
 - But α -Expansion has a guaranteed approximation quality (2-approx) and converges in a few iterations.

α -Expansion Move

- Basic idea:
 - Break multi-way cut computation into a sequence of binary s-t cuts.



α -Expansion Algorithm

1. Start with any initial solution
2. For each label “ α ” in any (e.g. random) order:
 1. Compute optimal α -expansion move (s-t graph cuts).
 2. Decline the move if there is no energy decrease.
3. Stop when no expansion move would decrease energy.

Example: Stereo Vision



Depth map

Original pair of “stereo” images

α -Expansion Moves

- In each α -expansion a given label “ α ” grabs space from other labels



initial solution

● -expansion

● -expansion

● -expansion

● -expansion

● -expansion

● -expansion

● -expansion


For each move, we choose the expansion that gives the largest decrease in the energy: \Rightarrow binary optimization problem

Topics of This Lecture

- Segmentation as Energy Minimization
 - Markov Random Fields
 - Energy formulation
- Graph cuts for image segmentation
 - Basic idea
 - s-t Mincut algorithm
 - Extension to non-binary case
- **Applications**
 - **Interactive segmentation**

GraphCut Applications: “GrabCut”

- Interactive Image Segmentation [Boykov & Jolly, ICCV’01]
 - Rough region cues sufficient
 - Segmentation boundary can be extracted from edges
- Procedure
 - User marks foreground and background regions with a brush.
 - This is used to create an initial segmentation which can then be corrected by additional brush strokes.



User segmentation cues




Additional
segmentation
cues





GrabCut: Data Model

Foreground
color



Background
color




Global optimum of
the energy

- Obtained from interactive user input
 - User marks foreground and background regions with a brush
 - Alternatively, user can specify a bounding box


GrabCut: Coherence Model

- An object is a coherent set of pixels:

$$\psi(x, y) = \gamma \sum_{(m,n) \in C} \delta[x_n \neq x_m] e^{-\beta \|y_m - y_n\|^2}$$




How to choose γ ?




Iterated Graph Cuts



Result



Color model
(Mixture of Gaussians)




Energy after
each iteration

GrabCut: Example Results



- ***This is included in the newest version of MS Office!***

Applications: Interactive 3D Segmentation



Summary: Graph Cuts Segmentation

- Pros

- Powerful technique, based on probabilistic model (MRF).
- Applicable for a wide range of problems.
- Very efficient algorithms available for vision problems.
- Becoming a de-facto standard for many segmentation tasks.

- Cons/Issues

- Graph cuts can only solve a limited class of models
 - Submodular energy functions
 - Can capture only part of the expressiveness of MRFs
- Only approximate algorithms available for multi-label case

References and Further Reading

- A gentle introduction to Graph Cuts can be found in the following paper:
 - Y. Boykov, O. Veksler, [Graph Cuts in Vision and Graphics: Theories and Applications](#). In *Handbook of Mathematical Models in Computer Vision*, edited by N. Paragios, Y. Chen and O. Faugeras, Springer, 2006.
- Read how the interactive segmentation is realized in MS Office 2010
 - C. Rother, V. Kolmogorov, Y. Boykov, A. Blake, [Interactive Foreground Extraction using Graph Cut](#), Microsoft Research Tech Report MSR-TR-2011-46, March 2011
- Try the GraphCut implementation at <http://www.cs.ucl.ac.uk/staff/V.Kolmogorov/software.html>