

Implementation of Databases

Assignment #3

Due on November 29, 2016

Prof. Dr. rer. pol. Matthias Jarke

Dr. rer. nat. Christoph Quix

Submitted by:

Sanchit Alekh, Idil Esen Zulfikar

MSc. Software Systems Engineering

Problem 1

Part 1

Let R be a relation of attributes A, B and C. Given is a query in relational algebra (RA).

- (a) Translate the RA query into an equivalent query in Domain Relational Calculus (DRC).

Solution

$$\begin{aligned} < A_1, B_1, C_1 > | \exists A_2, B_2, C_2, A_3, B_3, C_3 \\ & R_3(A_3, B_3, C_3), R_2(A_2, B_2, C_2), R_1(A_1, B_1, C_1) \\ & \wedge (A_2 = A_3) \wedge (B_1 = B_2) \wedge (B_2 = 1) \end{aligned}$$

- (b) Construct an equivalent tableau of the query.

Solution

a_1	b_1	c_1	
a_1	b_1	c_1	R
a_2	1	c_2	R
a_2	b_3	c_3	R
$b_1 = 1$			

Table 1: Tableau Representation of the Given Relational Algebra Query

Part 2

Given the following two tableau queries T1 and T2, decide whether they are equivalent, subsumed by each other, or not related at all.

Solution

Checking for $T_1 \subseteq T_2$, looking for mappings h from T_2 to T_1

$$h(a_5) = a_1$$

$$h(b) = b$$

$$h(c_4) = c_1$$

$$h(a) = a$$

$$h(b_5) = b_1$$

$$h(c_4) = c_1$$

Since a mapping h exists for each element $e \in T_2$, $T_1 \subseteq T_2$

Checking for $T_2 \subseteq T_1$, looking for mappings h' from T_1 to T_2

$$h'(a) = a$$

$$h'(b_1) = b_5$$

$$h'(c_1) = c_4$$

$$h'(a) = a$$

$$h'(b_3) = b_5$$

$$h'(c_3) = c_4$$

$$h'(a_1) = a_5$$

$$h'(b) = b$$

$$h'(c_1) = c_4$$

$$h'(a_1) = a_5$$

$$h'(b_2) = b$$

$$h'(c_2) = c_4$$

Since a mapping h' exists for each element $e' \in T_1$, $T_2 \subseteq T_1$

We have proved that $T_1 \subseteq T_2$ and $T_2 \subseteq T_1$

Therefore T_1 and T_2 are equivalent to each other.

Problem 2

Given is the following relational database schema (keys are underlined>):

1. TRC:

Specify the following query in the tuple relational calculus (TRC) and draw the corresponding quant graph: Find the names of the chefs, who created the brand, and the shops of his brand are located in the same city where the chef was born.

Solution

$$c.ChefName \mid c \in Chef \wedge \exists cr \in Create$$

$$cr.CID = c.CID \wedge$$

$$\exists s \in Shop \wedge s.BID = cr.BID \wedge s.City = c.CityOfBirth$$

The Quant Graph is shown in **Figure 1**.

2. Does the graph contain a cycle? What does the result mean for optimization?




Figure 1: Quant Graph representation of the Query

Solution

Yes, the quant graph in this question contains a cycle.

This means that the query implementation will be extremely sub-optimal no matter what join order is chosen. This is because the semi-joins give intermediate tuples which are much larger than the desired result. This makes it really difficult to employ any query optimization strategy.

Problem 3

Consider the following relational schema and SQL query. The schema captures information about employees, departments, and projects.

1. Compute the number of pages for each relation.

Solution

For Emp Relation:

- Total number of tuples = 20000
- Size of each tuple = 20 bytes
- Size of each page = 4000 bytes
- Hence, no. of pages = $\frac{\text{total size of all tuples}}{\text{size of one page}} = \frac{20000 \times 20}{4000} = 100$ pages

For DeptProj Relation:

- Total number of tuples = 5000
- Size of each tuple = 40 bytes
- Size of each page = 4000 bytes
- Hence, no. of pages = $\frac{\text{total size of all tuples}}{\text{size of one page}} = \frac{5000 \times 40}{4000} = 50$ pages

For Proj Relation:

- Total number of tuples = 1000
- Size of each tuple = 2000 bytes
- Size of each page = 4000 bytes
- Hence, no. of pages = $\frac{\text{total size of all tuples}}{\text{size of one page}} = \frac{2000 \times 1000}{4000} = 500$ pages

2. Consider the following query: Find all employees with age more than 30.” Assume that there is an unclustered index on age. Let the number of qualifying tuples be N. For what values of N is a sequential scan cheaper than using the index?

Solution

Cost of page-by-page access for the *Emp* relation = 100 I/Os, since there are 100 pages

Cost of I/Os using the unclustered index = Cost of index access + Page I/O for each qualifying tuple = 3 + N

Therefore, for sequential scan to be cheaper than using the index

$$3 + N > 100$$

$$\therefore N > 97$$

Therefore, $N > 97$ for sequential scan to be cheaper.

3. Consider the following query:

SELECT * FROM Emp E, DeptProj D WHERE E.did=D.did

- a Compute the costs for the query using a block-nested loop join.

Solution

For a block-nested loop join, the costs are given by:

$$\begin{aligned} & M + \left\lceil \frac{M}{B-2} \right\rceil \times N \\ &= 50 + \left\lceil \frac{50}{12-2} \right\rceil \times 100 = 50 + 500 = 550 \text{ I/Os} \end{aligned}$$

- b Suppose that there is a clustered hash index on *did* on *Emp*. Compute the costs for the query using an index-nested loop join.

Solution

If a clustered hash index is available on *Emp*, the best approach is to make it the inner relation and *DeptProj* the outer relation to use the Hash Index.

Cost of scanning the outer relation (*DeptProj*) = Number of pages = 50 I/O

Now, for each tuple in *DeptProj*, we have the following costs:

- Cost of Index Access = 3 I/Os
- Cost of retrieving the matching tuple = 1 I/O

Therefore, total cost:

$$M + t_M \times (\text{Cost of Index Access} + \text{Cost of Retrieving Matching Tuples}) \\ = 50 + 5000 \times (3 + 1) = 50 + 5000 \times 4 = 20050 \text{ I/Os}$$

- c Assume that both relations are sorted on the join column. Which join method should be applied and what are the costs?

Solution

Since both relations are sorted, the sort-merge join is appropriate. The cost of Sort-merge join is calculated as the cost of sorting the relations on the join column + merging them, i.e. $O(M \log M) + O(N \log N) + M + N$ I/Os. In this case, both relations are already sorted on the join column, so the only thing should be calculated is the merge operation, which equals $M + N$.

$$\therefore \text{The cost of joining} = M + N = 100 + 50 = 150 \text{ I/Os}$$

- d Suppose that there is a clustered B+ tree index on *did* on *Emp* and *DeptProj* is sorted on *did*. Which join method should be applied and what are the costs?

Solution

\therefore There is a clustered B+ tree index on *Emp* on the key *did*, the values in the *Emp* relation must be sorted on *did*.

Furthermore, *DeptProj* is also sorted on *did*.

\therefore The most efficient approach is to use the sort-merge join, since both the relations are already sorted on the merge column.

The only costs to consider, would be the costs for joining, i.e. $M + N$

$$\therefore \text{Total Cost for the join} = 100 + 50 = 150 \text{ I/Os.}$$